Calibrate

INVESTOR DECK · 2026

# Know what every AI dollar buys. *Then spend fewer of them.*

Calibrate is the FinOps platform for enterprise AI. We attribute, control, and cut LLM spend without losing quality.



**Will Marriott**
CEO

**Dundy Pasupuleti**
CTO

SAN FRANCISCO

01 — PROBLEM

# AI bills grow faster than anyone can explain them.

INVOICE MONTHLY

**Model API provider**

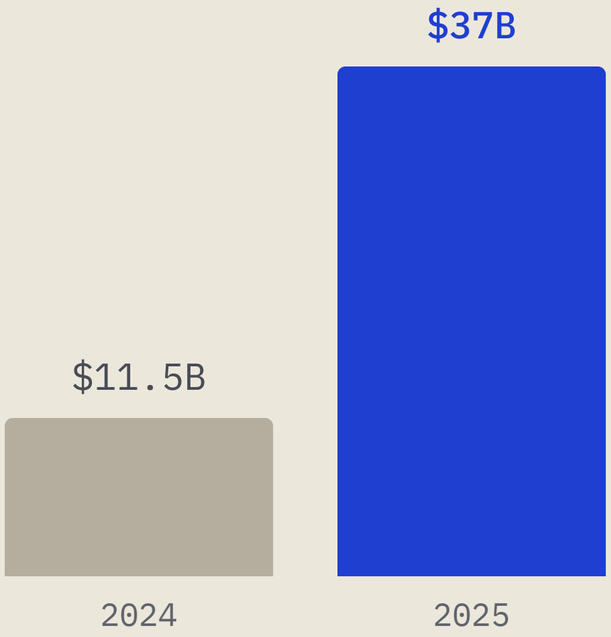
| | |
|---|---|
| LLM API usage | $ ??? |
| Line items | 1 |

**Total due** $ ???

| | |
|---|---|
| FINANCE ASKS | Which team spent it? |
| PRODUCT ASKS | Which feature drove it? |
| ENGINEERING ASKS | Did that call need the biggest model? |

*One invoice. No link to teams, features, or value.*

# Enterprises tripled generative AI spend in one year.



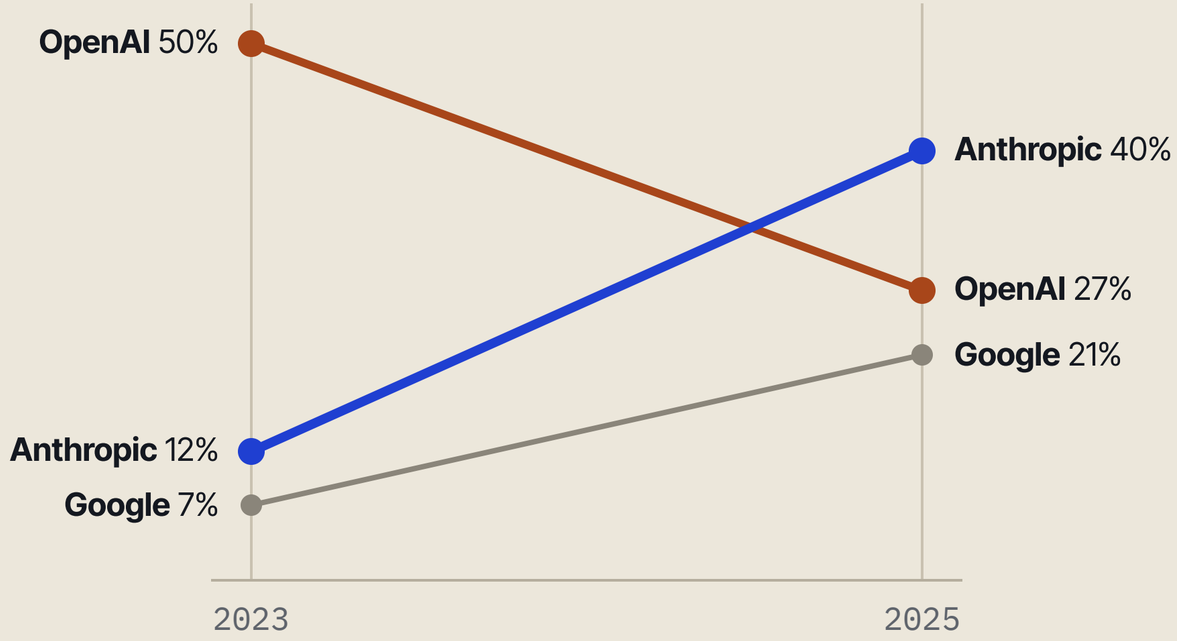

**$37B** — Enterprise generative AI spend in 2025

**3.2×** — Growth in a single year, up from $11.5B in 2024

**$18B** — Went to infrastructure: foundation models, training, and tooling

Source: Menlo Ventures, “2025: The State of Generative AI in the Enterprise,” Dec 2025

# The top model provider changed hands in two years.




## Spend follows quality, so it moves.

Three labs now hold 88% of enterprise LLM API spend, and the lead flips as models leapfrog each other. Every switch reprices the bill.

**Model choice is now a cost lever, not a one-time decision.**

Share of enterprise LLM API spend · Source: Menlo Ventures, Dec 2025

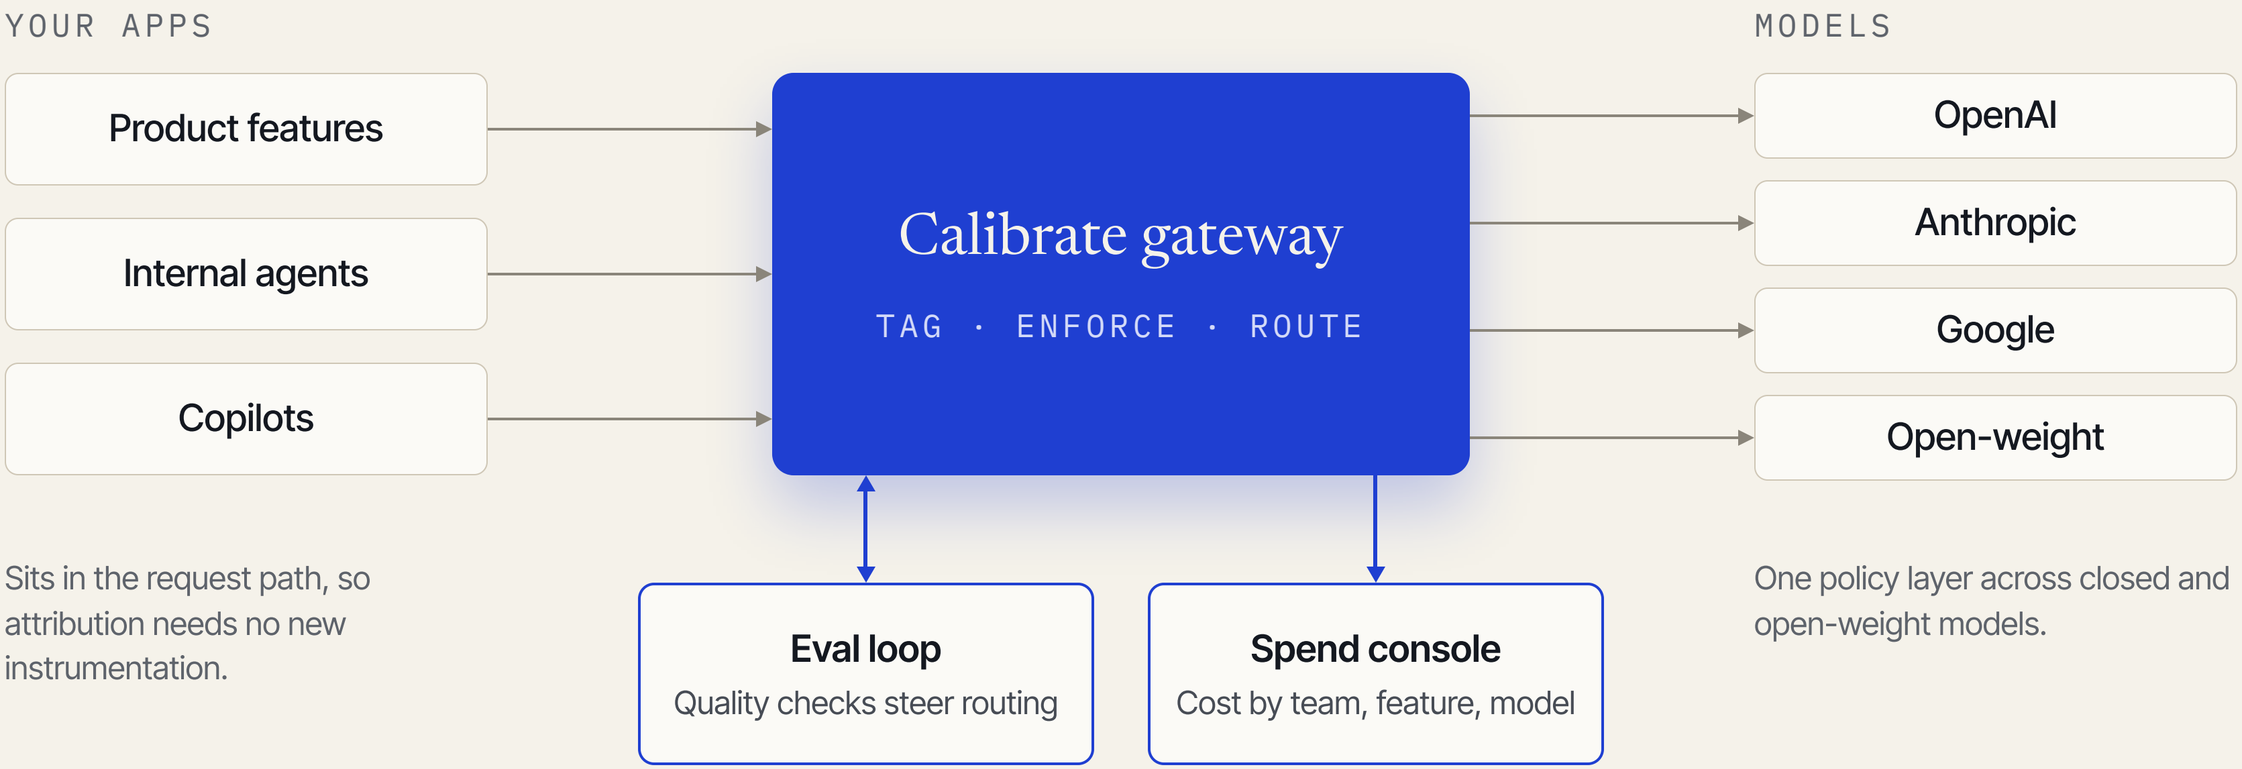
THE GAP

# Every enterprise now runs a portfolio of models. *Nobody manages it like one.*

Cloud spend got FinOps. AI spend still gets a single line on an invoice.

03 — HOW IT WORKS

# One gateway in front of every model call.




Sits in the request path, so attribution needs no new instrumentation.

One policy layer across closed and open-weight models.

03 — SOLUTION

# Calibrate sees, governs, and cuts every LLM dollar.

01

## Attribute

*Who spent it, and what did it buy?*

- By team, feature, and customer
- Cost per outcome, not per token
- Flags spend with no return

02

## Control

*What are we allowed to spend?*

- Approved models per team
- Budgets, caps, and alerts
- An audit trail on every call

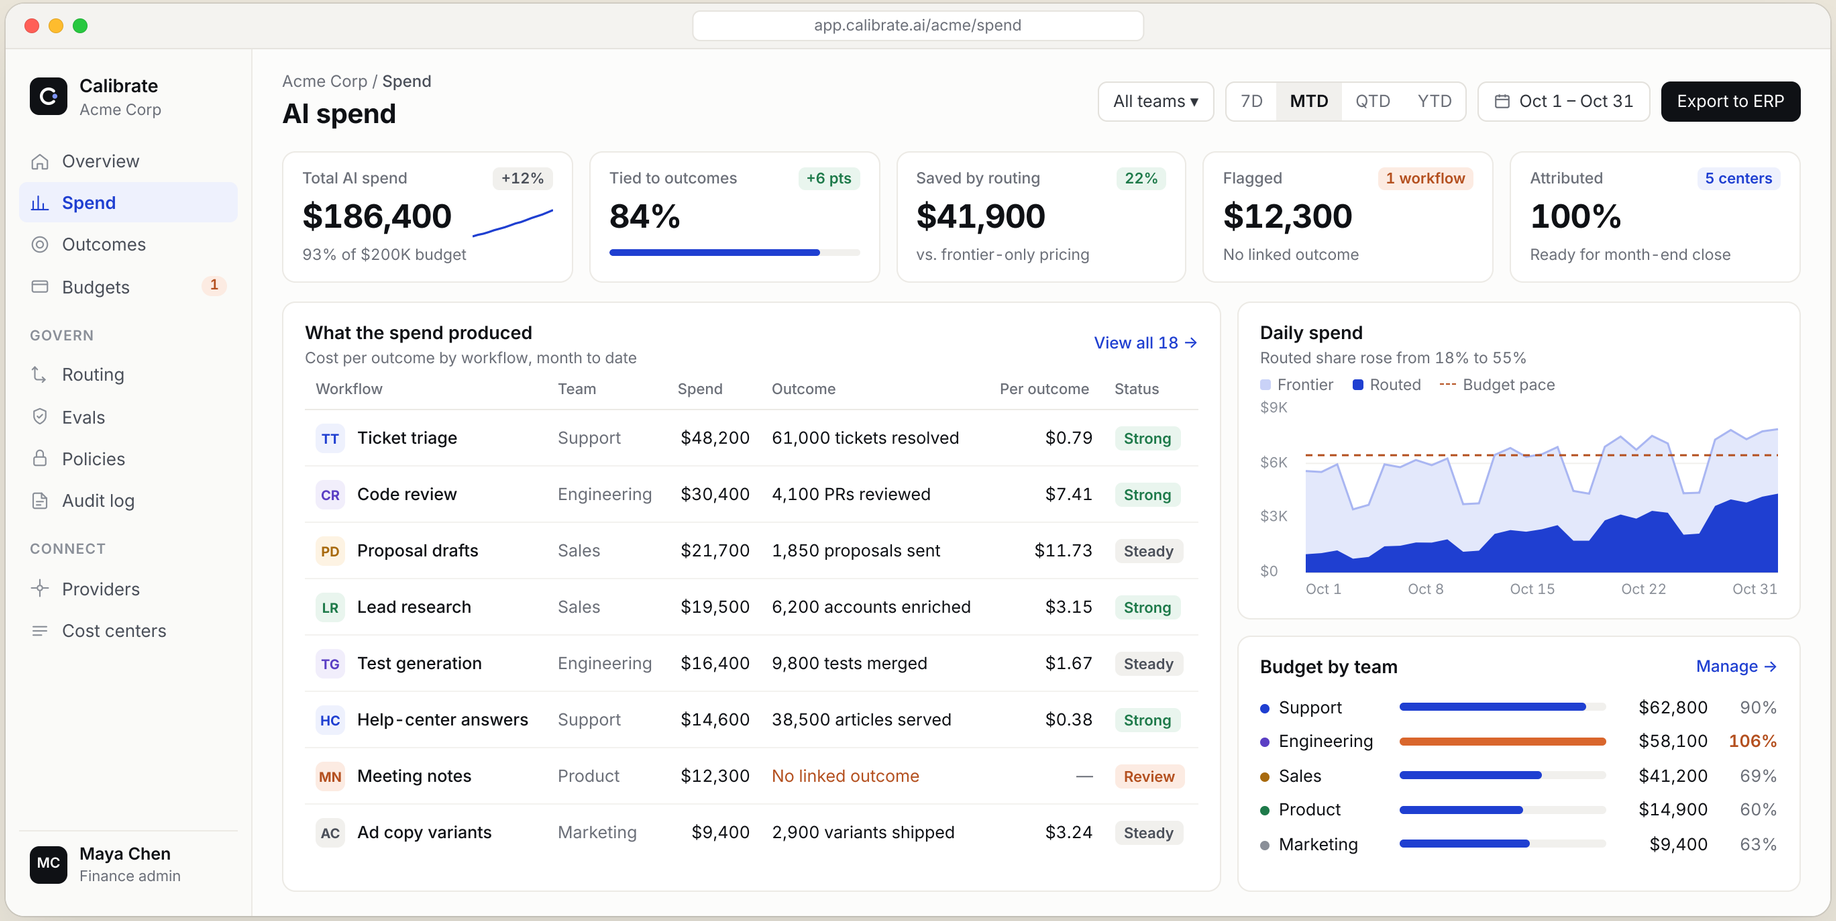
03

## Optimize

*What should this cost?*

- Cheapest model that passes evals
- Automatic, with override
- Open, closed, or your own models

# Every AI dollar has an owner and an outcome.



03 — ENTERPRISE CONTROLS

# Controls on every call, from day one.

ACCESS

## Who can use what

Which models, tools, and data each team can reach.

- SSO and SCIM
- Role-based permissions
- Approved models per team

RULES AND CAPS

## Under which rules

Policies enforced before the money is spent.

- Budgets, hard caps, and alerts
- Blocked data and actions
- Cheaper fallback at a cap

AUDIT

## On the record

Every call logged and ready for review by security or finance.

- Who called what, and when
- Model and version per call
- Exports for compliance and close

Security signs off on day one. *Finance closes the month from the same record.*

04 — WHY WE WIN

# Built for the CFO, not just the engineer.

| CATEGORY | EXAMPLES | BUILT FOR | COST IS… |
|---|---|---|---|
| **Agent platforms** | Guild.ai | CTO, CISO | One panel in a control plane |
| **Eval + gateway** | Maxim, Portkey, Helicone | AI engineers | A metric beside quality and uptime |
| **Model routers** | Not Diamond, Martian, OpenRouter | Developers | Decided one request at a time |
| **Cloud FinOps** | Vantage, CloudZero | Finance | Read from the invoice, after the fact |
| **Calibrate** | Spend, policy, and quality in one loop | **CFO + AI platform** | **The product itself** |

*Engineering platforms show cost. Calibrate turns it into chargebacks, budgets, and savings that evals prove.*

04 — BUSINESS MODEL

# Our revenue is tied to what we save you.

WHAT WE CHARGE

## A share of the savings

Priced on verified savings against the customer's frontier-only baseline.

WHAT WE DON'T TOUCH

## Model spend

Passed straight through to the provider. No markup on tokens.

ILLUSTRATIVE MATH

**$1M**

a year in model API spend

**$100K**

saved when 10% of spend routes to cheaper models that pass evals

**Calibrate** a share of the savings

**Customer** keeps the rest

## We make money when the bill shrinks, *not when it grows.*

Savings are verified by the customer's own evals. We'll set the exact split with our first pilot customers.

04 — GO-TO-MARKET

# Land with the AI platform team. Expand to finance.

01 · LAND

## Attribute

The platform team puts the gateway in front of existing calls and sees cost by team and feature.

02 · EXPAND

## Control

Finance sets budgets and chargebacks on the same data, with no new integration.

03 · GROW

## Optimize

Eval-gated routing turns visibility into savings, and savings into renewals.

CHAMPION

**Head of AI platform**

ECONOMIC BUYER

**CFO or FinOps lead**

FIRST TARGETS

**AI teams paying more than one provider**

05 — TEAM

# Finance meets production AI infrastructure.

WM

## Will Marriott

CEO

General Atlantic · Goldman Sachs · Accel-KKR

Investor on General Atlantic's technology team ($100B+ AUM). Vets companies and markets, and serves as a board observer on portfolio AI automation rollouts.

Previously a Goldman Sachs investment banker, advising on $5B+ of M&A and financings across software, defense, and media.

DP

## Dundy Pasupuleti

CTO

Pilot · Apple · Santa Clara University

Senior Software Engineer, AI at Pilot, running production LLM agents on AWS. At Apple, built fused CUDA kernels and TensorRT-LLM integrations; also built model quantization pipelines.

M.S. in CS. CKA. Open-source contributor to vLLM, NVIDIA Dynamo, FlashInfer, Ray Serve, NVIDIA GPU Operator, and Kueue.

05 — STATUS AND ASK

# Early, honest, and moving fast.

| | |
|---|---|
| MVP | Building now |
| CUSTOMERS | In talks with clients for pilot runs |

## Know what every AI dollar buys. *Then spend fewer of them.*

**wmarriott99@gmail.com** · **dundysm@gmail.com**

THE ASK

## Join Y Combinator.

To ship the MVP and turn pilot talks into live pilots with design partners.

WHAT IT UNLOCKS

01 Ship the gateway and spend console

02 Sign design partners

03 Publish our first routed-savings case study